

# Kafka Consumer

Implementation details, covering setup, your Kafka-based consumption service, and other pertinent aspects

There are several reasons for having the decoupled DB (producer-consumer) in the first place:

1. Performance. If a bunch of consumers query the producer DB, it can stress the database, slowing down the producer service. Each consumer would independently hit the producer DB for the same data, multiplying the load.
2. Coupling. Any schema change, maintenance, or downtime in the producer DB directly impacts all consumers.
3. Replay. Once the consumer reads the data, if it misses it, there's no built-in way to "replay" the event without Kafka
4. Transform. Often, data needs to be extracted, transformed, and only the relevant attributes must be loaded and used by consumers.

The main purpose of the consumer to be developed is to "consume" the images from the producer DB to the consumer DB, ensuring the data between the databases is synchronized. This shall happen in real time, so as a background process or daemon.

## Consumer Job

To achieve that, I made use of the kafka-python and pymongo libraries. Consumer Job first creates a connection to the Mongo database and the Kafka bootstrap servers:

```
1 # Lifespan
2 def main():
3     # Mongo
4     logger.info("Connecting to Mongo")
5     db = get_db()
6
7     # Kafka
8     logger.info("Connecting to Kafka")
9     consumer = KafkaConsumer(
10         *TOPICS,
11         bootstrap_servers=KAFKA,
12         group_id=GROUP_ID,
13         auto_offset_reset="earliest", # when there's no committed offset yet
14         enable_auto_commit=False,
15         value_deserializer=lambda x: json.loads(x.decode("utf-8")),
16     )
17
18     logger.info("Kafka consumer started")
19     consume_messages(consumer, db)
```

I define the topics I want to subscribe to, the consumer group ID, and disable auto-committing. Importantly, I set `auto\_offset\_reset` to earliest to start reading from the beginning of the

topic if there is no committed offset yet, and deserialize the received payload to a Python dictionary object for easier access later.

Then it blocks in `for m in consumer` waiting for Kafka messages to arrive. If no messages are available, the broker waits up to `fetch\_max\_wait\_ms` (default 500 ms) to see if new messages arrive. For each received message, it upserts the value in the consumer database, ensuring it stays up-to-date with the producer database. And since we have soft delete (changing the flag), this is enough. Consumer Job also logs the changes it made.

```
● ○ ●
1 # Consumer loop
2 def consume_messages(consumer: KafkaConsumer, db: Database):
3     for m in consumer:
4         try:
5             value, topic = m.value, m.topic
6
7             collection = get_collection_by_type(db, topic)
8
9             result = collection.update_one(
10                 {"id": value["id"]}, {"$set": value}, upsert=True
11             )
12
13             if result.upserted_id:
14                 logger.info(f"● Inserted {topic} #{value['id']}")
15             elif result.modified_count > 0:
16                 logger.info(f"● Updated {topic} #{value['id']}")
17             else:
18                 logger.info(f"● No changes {topic} #{value['id']}")
19         except Exception as e:
20             logger.error(f"Error processing message: {e}")
21         finally:
22             consumer.commit()
23
```

```
INFO Kafka.Consumer.Fetcher Resetting offset
INFO root ● Inserted potato #2153
INFO root ● Inserted pepper #2475
INFO root ○ Updated potato #1723
INFO root ● Inserted potato #2154
INFO root ● Inserted potato #2155
INFO root ● Inserted potato #2156
INFO root ● Inserted potato #2157
INFO root ● Inserted potato #2158
INFO root ● Inserted potato #2159
INFO root ● Inserted tomato #16011
INFO root ○ Updated pepper #1615
INFO root ○ Updated tomato #9415
INFO root ○ Updated tomato #11038
INFO root ● Inserted tomato #16012
```

## Error handling

### During looping

In case an error is encountered parsing the message to value and topic, or e.g., connection to the database is dropped, the message is dropped for simplicity (with `finally: commit()`).

One solution to this would be to move the commitment to the successful execution only (`try` part) - this way, Consumer Job will refetch the message from the last offset and retry until success. But if the problem is in the message itself, it becomes a dead letter, stopping the synchronisation job indefinitely. A compromise would be to set the retry backoff to N times or send failed messages to a separate dead-letter topic.

I tried killing the Kafka broker while looping for messages as well. Kafka consumer enters reconnection loop without throwing an exception and exiting until it can successfully reach the broker.

### During creation

I deliberately decided not to introduce `/ready` and `/live` probes to the job and not to silence the critical errors. So if any error is thrown in `main()`, the application will exit with an error and Kubernetes will restart it.

## Sync vs Async

Consumer shall also expose two endpoints:

`/image-plant/{image\_type}/{image\_id}` and `/image-plant/{image\_type}/total`. First, I put it all in one FastAPI app, making the consumer job run as a background process, and the foreground was serving api requests. This required connections to Kafka and the MongoDB, made asynchronous to yield control from the job to the api when waiting for new messages to arrive or making changes to the database. But shortly after, I realized that this is bad design in the first place to have API and Job in a single app, so I decided to decouple them into two microservices, namely Consumer Api and Consumer Job. This also removed the need to make Kafka and MongoDB async because in Consumer Job thread now has nothing else to do, so there is no need to yield control to anything.

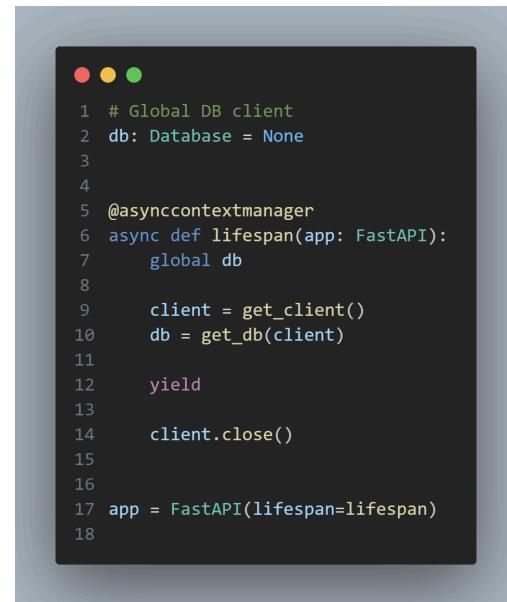
## Consumer API

Consumer API is a simple FastAPI app implementing the required endpoints, but also introducing /ready and /live probes, which are essential for Kubernetes to understand whether the API is not stuck.

A screenshot of a terminal window displaying a Python script for a FastAPI application. The script includes imports for FastAPI, Pydantic, and MongoDB operations. It defines a `readiness` endpoint that returns "ready", a `liveness` endpoint that returns "alive", and two main endpoints: `/image-plant/{image\_type}/total` which returns the count of documents in a MongoDB collection, and `/image-plant/{image\_type}/{image\_id}` which finds a document by ID and returns its details. Error handling is provided for non-existent documents.

```
1 # Probes (will hang in case something bad happened)
2 @app.get("/ready")
3 async def readiness():
4     return "ready"
5
6
7 @app.get("/live")
8 async def liveness():
9     return "alive"
10
11
12 # endpoints
13 @app.get("/image-plant/{image_type}/total")
14 def get_total_images(image_type: ImageType):
15     collection = get_collection_by_type(db, image_type.value)
16
17     total = collection.count_documents({})
18     return total
19
20
21 @app.get("/image-plant/{image_type}/{image_id}")
22 def get_image_by_id(
23     image_type: ImageType,
24     image_id: int,
25 ):
26     collection = get_collection_by_type(db, image_type.value)
27
28     doc = collection.find_one({"id": image_id})
29     if not doc:
30         raise HTTPException(status_code=404, detail="Image not found")
31
32     doc.pop("_id", None)
33     return doc
34
```

It also defines a global client for interaction with the consumer database.



```
1 # Global DB client
2 db: Database = None
3
4
5 @asynccontextmanager
6 async def lifespan(app: FastAPI):
7     global db
8
9     client = get_client()
10    db = get_db(client)
11
12    yield
13
14    client.close()
15
16
17 app = FastAPI(lifespan=lifespan)
18
```

## Sync vs Async

FastAPI allows for async endpoints. However, since Job uses a synchronous MongoDB client, rather than rewriting everything to the async Motor library, I decided to keep the common MongoDB client connection code synchronous and maintain synchronous endpoints. This way, Uvicorn threadpool handles parallel requests instead of async (green) threads. From the docs: PyMongo is thread-safe and provides built-in connection pooling for threaded applications. So, PyMongo handles concurrent requests via connection pooling, so multiple threads can query the DB in parallel.

## Challenges

Main challenges here were to come up with this decoupled architecture, decide sync/async, and find the config names for Kafka topics, DB and collection names, and the format of the messages.

## Scalability

Both the Consumer API and the Consumer Job services have the resource requirements and limits defined, which allow them to burst in case of traffic increase.

## API

If the intensity of data flow increased, thead pools may still not be enough, and fully async API routes and mongo connections will still make a difference.

Resource requests and limits I started with:



This clearly showed that both limits and requests must have been decreased, and I decreased them to:

```

● ○ ●
1 resources:
2   requests:
3     cpu: "10m"
4     memory: "64Mi"
5   limits:
6     cpu: "50m"
7     memory: "128Mi"
8

```

## Job

Since the job processes messages one by one, there is not much I could change here. Async will not help here.

What could be done is to introduce Consumer Job replicas, each subscribed to one of the potato/tomato/pepper topics.

Or even more, topics could be divided into partitions, but ensuring images with specific ID always get to the same partition, and Consumer Job replicas will subscribe to those partitions. This way, there will be no race conditions between changes to the same image (because Kafka guarantees the order of messages within a single partition).

To make the number of Consumer Job replicas adjust to the load (scale up and down like HPA), KEDA (Kubernetes Event-Driven Autoscaler) could be introduced, and a ScaledObject for Kafka consumer lag can be created. The advantage is that KEDA automatically scales the number of Kafka consumer pods up and down based on real Kafka lag. The disadvantage is that it has to be learned and configured.

Initial resource requests and limits turned out to be high as well:



So I decreased them to:



On GKE, I dedicate enough resources to the nodes so that Kubernetes pods can scale vertically to their limits in resources in case of CPU or memory load increase.

## Batch processing vs Stream processing

The first thing that can be noticed on the Grafana dashboard is that there is a bunch of images already present in the Producer DB and being picked up by the db-synchronizer job. While new images are being pushed to Kafka. This results in a significantly higher number of images processed by batching vs streaming.

The second thing is that I would expect the graph of the number of processed images by the batching to look like a step function with large jumps. However, on the first run (when both the

DB synchronizer and the consumer job were up), it rather resembled the performance of stream processing. That could be because of the large number of images to process, and even though it activates once every 10 minutes, it processes a batch of images one by one.



As expected, the number of images processed by the stream grows smoothly because as soon as an image arrives at the Kafka topic, it is picked up and consumed.

A similar observation can be made about the size of the images processed.

After turning the DB synchronizer off, the number of images and the size of the images processed by the batch stopped at the same level, which is expected.



Interestingly, after a restart DB synchronizer tried to process another batch and encountered an error, which stopped the processing at some number of images (without processing the whole batch). As scheduled, it tries again once 10 minutes, but reaches same error.

```

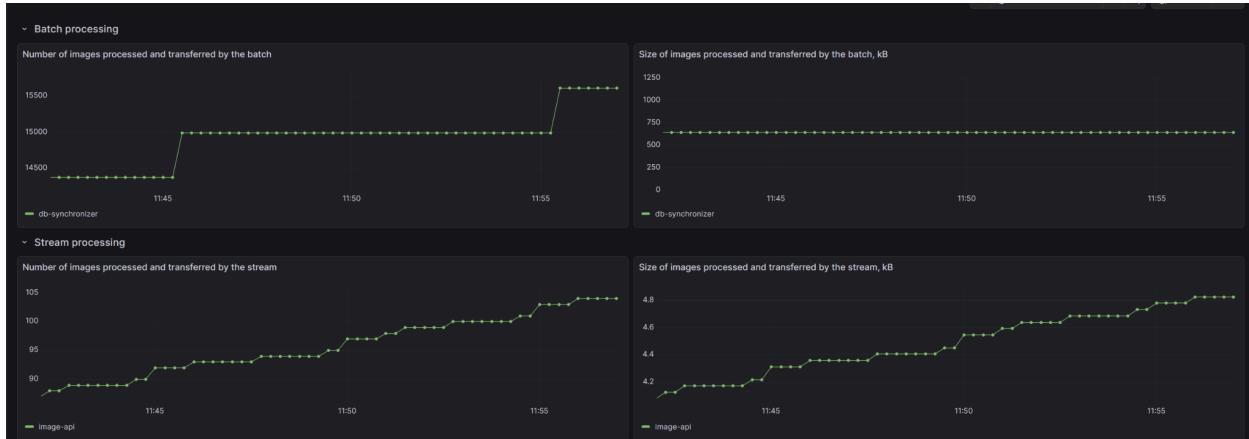
127.0.0.1 - - [19/Dec/2025 23:27:40] "GET / HTTP/1.1" 200 -
127.0.0.1 - - [19/Dec/2025 23:27:41] "GET /static/css/main.css HTTP/1.1" 304 -
127.0.0.1 - - [19/Dec/2025 23:27:45] "POST /synchronize-leaf-image-dbs-job-start HTTP/1.1" 200 -
synchronize_leaf_image_dbs_job: Started!
synchronize_leaf_image_dbs_job: Database is deployed. Starting to synchronize the databases.
127.0.0.1 - - [19/Dec/2025 23:27:45] "GET /static/css/main.css HTTP/1.1" 304 -
Started to work with batch: 1 of plant: potato with batchsize: 50 and from image_id: 1 out of 2360
Started to work with batch: 2 of plant: potato with batchsize: 50 and from image_id: 51 out of 2360
Job "synchronize_leaf_image_dbs_job (trigger: interval[0:10:00], next run at: 2025-12-10 23:37:45 UTC)" raised an exception
Traceback (most recent call last):
  File "/usr/local/lib/python3.10/site-packages/apscheduler/executors/base.py", line 125, in run_job
    retval = job.func(*job.args, **job.kwargs)
  File "/app/db_synchronizer_job.py", line 97, in synchronize_leaf_image_dbs_job
    elif image_db["updated_at"] is None or image_db["updated_at"] < image.updated_at:
TypeError: '<' not supported between instances of 'str' and 'datetime.datetime'
TypeError: 'None' object has no attribute 'updated_at'

Containers / k8s_db-synchronizer_db-synchronizer-76bfdf8d58-m9hq_leaf-image-management-system_e9082544-ce80-492b-825f-8628d1b29544_0
  k8s_db-synchronizer_db-synchronizer-76bfdf8d58-m9hq_leaf-image-management-
  system_e9082544-ce80-492b-825f-8628d1b29544_0
  STATUS Running (4 hours ago) ⚡ 🔍 🗑
  @a67874cac8b ⚡ 1221994/db_synchronizer_job-latest

Logs Inspect Bind mounts Exec Files Stats
Started to work with batch: 4 of plant: potato with batchsize: 50 and from image_id: 151 out of 2292
Started to work with batch: 5 of plant: potato with batchsize: 50 and from image_id: 201 out of 2292
Job "synchronize_leaf_image_dbs_job (trigger: interval[0:10:00], next run at: 2025-12-10 20:21:54 UTC)" raised an exception
Traceback (most recent call last):
  File "/usr/local/lib/python3.10/site-packages/apscheduler/executors/base.py", line 125, in run_job
    retval = job.func(*job.args, **job.kwargs)
  File "/app/db_synchronizer_job.py", line 97, in synchronize_leaf_image_dbs_job
    elif image_db["updated_at"] is None or image_db["updated_at"] < image.updated_at:
TypeError: '<' not supported between instances of 'str' and 'datetime.datetime'
Run time of job "synchronize_leaf_image_dbs_job (trigger: interval[0:10:00], next run at: 2025-12-10 20:21:54 UTC)" was missed by 0:00:04.179907
Run time of job "synchronize_leaf_image_dbs_job (trigger: interval[0:10:00], next run at: 2025-12-10 20:41:54 UTC)" was missed by 0:00:04.126654
Run time of job "synchronize_leaf_image_dbs_job (trigger: interval[0:10:00], next run at: 2025-12-10 21:01:54 UTC)" was missed by 0:00:04.081777
Run time of job "synchronize_leaf_image_dbs_job (trigger: interval[0:10:00], next run at: 2025-12-10 21:21:54 UTC)" was missed by 0:00:04.081447
Run time of job "synchronize_leaf_image_dbs_job (trigger: interval[0:10:00], next run at: 2025-12-10 21:41:54 UTC)" was missed by 0:00:04.082168
Run time of job "synchronize_leaf_image_dbs_job (trigger: interval[0:10:00], next run at: 2025-12-10 21:51:54 UTC)" was missed by 0:00:04.0816941
Run time of job "synchronize_leaf_image_dbs_job (trigger: interval[0:10:00], next run at: 2025-12-10 21:31:54 UTC)" was missed by 0:00:04.761769
Run time of job "synchronize_leaf_image_dbs_job (trigger: interval[0:10:00], next run at: 2025-12-10 21:41:54 UTC)" was missed by 0:00:04.264793
Run time of job "synchronize_leaf_image_dbs_job (trigger: interval[0:10:00], next run at: 2025-12-10 21:51:54 UTC)" was missed by 0:00:04.899324
Run time of job "synchronize_leaf_image_dbs_job (trigger: interval[0:10:00], next run at: 2025-12-10 22:01:54 UTC)" was missed by 0:00:04.858797
Run time of job "synchronize_leaf_image_dbs_job (trigger: interval[0:10:00], next run at: 2025-12-10 22:11:54 UTC)" was missed by 0:00:04.850713
Run time of job "synchronize_leaf_image_dbs_job (trigger: interval[0:10:00], next run at: 2025-12-10 22:31:54 UTC)" was missed by 0:00:04.843503
Run time of job "synchronize_leaf_image_dbs_job (trigger: interval[0:10:00], next run at: 2025-12-10 22:41:54 UTC)" was missed by 0:00:04.833379
Run time of job "synchronize_leaf_image_dbs_job (trigger: interval[0:10:00], next run at: 2025-12-10 22:51:54 UTC)" was missed by 0:00:04.833334
Run time of job "synchronize_leaf_image_dbs_job (trigger: interval[0:10:00], next run at: 2025-12-10 23:01:54 UTC)" was missed by 0:00:04.550338
Run time of job "synchronize_leaf_image_dbs_job (trigger: interval[0:10:00], next run at: 2025-12-10 23:11:54 UTC)" was missed by 0:00:04.766296
Run time of job "synchronize_leaf_image_dbs_job (trigger: interval[0:10:00], next run at: 2025-12-10 23:21:54 UTC)" was missed by 0:00:04.676408
Run time of job "synchronize_leaf_image_dbs_job (trigger: interval[0:10:00], next run at: 2025-12-10 23:31:54 UTC)" was missed by 0:00:04.633825
Run time of job "synchronize_leaf_image_dbs_job (trigger: interval[0:10:00], next run at: 2025-12-10 23:41:54 UTC)" was missed by 0:00:04.653179

```

However, now the Grafana panel showed more like a step function with jumps (of the number of images it had time to process before an error), which confirms previous assumption.



One can improve a batch approach to behave almost like stream processing by using micro-batches, triggering batches as soon as new data arrives (or in short time intervals) instead of on a fixed schedule - in fact, this is what Kafka Consumer stream processing does internally with `fetch\_min\_bytes` and `fetch\_max\_wait\_ms` parameters. Stream processing is considered better because it provides:

- Real-time, low-latency processing
- Continuous, stable resource usage
- Incremental computation instead of heavy periodic jobs
- Built-in fault tolerance and checkpointing
- Immediate reactions to events (analytics, alerts, automation)

This makes it generally more responsive, efficient, and scalable than traditional batch systems.

Also note that to scale batch processing, one could split data into chunks and fit them into horizontally scaled DB synchronizers, based on a custom consumption lag.

## K8s Resources

I introduced Deployment objects for the Consumer API and Consumer Job. Both have resources defined, possibility to change config of the apps with env (like MONGO DB URL or KAFKA BOOTSTRAP SERVERS URL), and specified /ready and /live probes for API.

I also put ClusterIP services in front of the Consumer API and the Consumer Job. And since the API should be made accessible, I also configured Ingress to forward connections to the Consumer API endpoint. So the only entry point to the consumer is through ingress.

## Deployment | GKE

I developed the Consumer locally but tested it on GCP as well.

Moving to GCP, I had to adjust the resource requests/limits (decreasing them), which still works fine locally, so I don't have separate manifests for staging and production. The only difference between staging and production is the presence of an ingress object - it is not necessary for local development since I configured port forwarding to the Consumer API service in [open.stg.sh](#), and if one were to use ingress locally, them will need to start an ingress controller.

As for cluster, I decided to use 3 nodes instead of 1 to make the system more resilient, and this is recommended by GCP.

### vCPU / Memory

I started with lower values, but successfully deployed with 4 vCPU and 4 GB per node.

An equivalent code snippet of my cluster creation:

```
gcloud beta container clusters create cluster-1 \
    --project cloud-computing-476715 \
    --zone us-east1-b \
    --no-enable-basic-auth \
    --cluster-version 1.33.5-gke.1308000 \
    --release-channel regular \
    --machine-type e2-custom-4-4096 \
    --image-type COS_CONTAINERD \
    --disk-type pd-standard \
    --disk-size 40 \
```

Node pools [Create user-managed node pool](#)

[Filter](#) Filter node pools

Name	Status	Version	Number of nodes	Machine type	Image type	Autoscaling	Default IPv4 Pod IP address range	⋮
pool-1	Ok	1.33.5-gke.1308000	3	e2-custom-4-4096	Container-Optimized OS with containerd (cos_containerd)	Off	10.8.0.0/14	

Nodes

[Filter](#) Node type : User-managed [Filter nodes](#)

Name	Status	Node type	CPU requested	CPU allocatable	Memory requested	Memory allocatable	Storage requested
gke-cluster-1-pool-1-5f572885-0zkr	Ready	User-managed	1.73 CPU	3.92 CPU	2.91 GB	2.94 GB	0 B
gke-cluster-1-pool-1-5f572885-rts9	Ready	User-managed	941 mCPU	3.92 CPU	2.6 GB	2.94 GB	0 B
gke-cluster-1-pool-1-5f572885-wgfh	Ready	User-managed	1.85 CPU	3.92 CPU	2.66 GB	2.94 GB	0 B

## Workloads

Workloads						
Cluster		Namespace				
cluster-1 (us-east1-b)		leaf-image-managemen...		Reset	Save	
<a href="#">Overview</a> <a href="#">Observability</a> <a href="#">Cost Optimization</a>						
<input type="checkbox"/> <a href="#">Name ↑</a>	Status	Type	Pods	Node type	Namespace	Cluster
<input type="checkbox"/> <a href="#">camera</a>	<span>✓ OK</span>	Deployment	1/1	User-managed	leaf-image-management-system	<a href="#">cluster-1</a>
<input type="checkbox"/> <a href="#">consumer-api</a>	<span>✓ OK</span>	Deployment	1/1	User-managed	leaf-image-management-system	<a href="#">cluster-1</a>
<input type="checkbox"/> <a href="#">consumer-job</a>	<span>✓ OK</span>	Deployment	1/1	User-managed	leaf-image-management-system	<a href="#">cluster-1</a>
<input type="checkbox"/> <a href="#">db-synchronizer</a>	<span>✓ OK</span>	Deployment	1/1	User-managed	leaf-image-management-system	<a href="#">cluster-1</a>
<input type="checkbox"/> <a href="#">image-api</a>	<span>✓ OK</span>	Deployment	1/1	User-managed	leaf-image-management-system	<a href="#">cluster-1</a>
<input type="checkbox"/> <a href="#">kafka-broker</a>	<span>✓ OK</span>	Deployment	1/1	User-managed	leaf-image-management-system	<a href="#">cluster-1</a>
<input type="checkbox"/> <a href="#">mongodb-consumer</a>	<span>✓ OK</span>	Deployment	1/1	User-managed	leaf-image-management-system	<a href="#">cluster-1</a>
<input type="checkbox"/> <a href="#">mongodb-producer</a>	<span>✓ OK</span>	Deployment	1/1	User-managed	leaf-image-management-system	<a href="#">cluster-1</a>
<input type="checkbox"/> <a href="#">users</a>	<span>✓ OK</span>	Deployment	1/1	User-managed	leaf-image-management-system	<a href="#">cluster-1</a>
<input type="checkbox"/> <a href="#">zookeeper</a>	<span>✓ OK</span>	Deployment	1/1	User-managed	leaf-image-management-system	<a href="#">cluster-1</a>

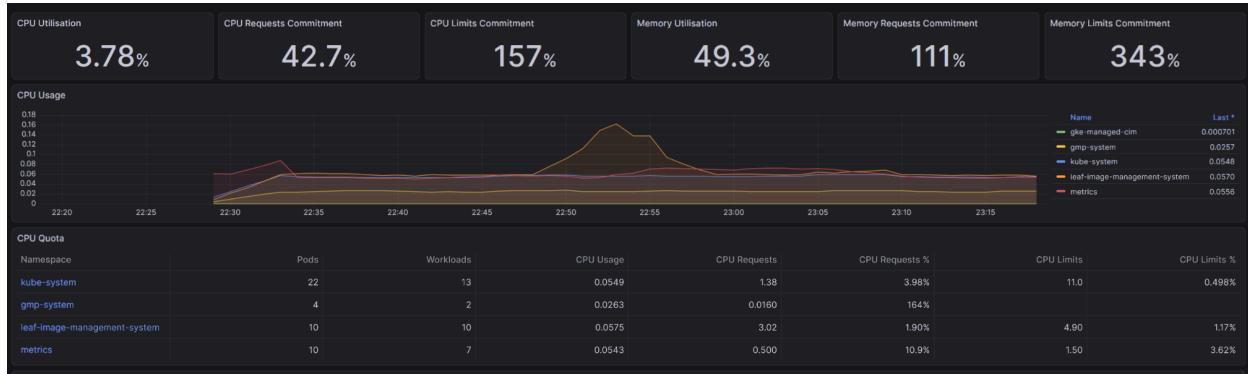
## Services

Services						
<input type="checkbox"/> <a href="#">Name ↑</a>	Status	Type	Endpoints	Pods	Namespace	Clusters
<input type="checkbox"/> <a href="#">camera</a>	<span>✓ OK</span>	Node Port	34.118.235.114:5050 TCP	1/1	leaf-image-management-system	<a href="#">cluster-1</a>
<input type="checkbox"/> <a href="#">consumer-api</a>	<span>✓ OK</span>	Cluster IP	34.118.229.243	1/1	leaf-image-management-system	<a href="#">cluster-1</a>
<input type="checkbox"/> <a href="#">consumer-job</a>	<span>✓ OK</span>	Cluster IP	34.118.227.127	1/1	leaf-image-management-system	<a href="#">cluster-1</a>
<input type="checkbox"/> <a href="#">db-synchronizer</a>	<span>✓ OK</span>	Node Port	34.118.228.109:5050 TCP	34.118.228.109:8050 TCP	1/1	leaf-image-management-system
<input type="checkbox"/> <a href="#">image-api</a>	<span>✓ OK</span>	Node Port	34.118.235.87:8080 TCP	34.118.235.87:8050 TCP	1/1	leaf-image-management-system
<input type="checkbox"/> <a href="#">kafka-service</a>	<span>✓ OK</span>	Cluster IP	None	1/1	leaf-image-management-system	<a href="#">cluster-1</a>
<input type="checkbox"/> <a href="#">mongodb-consumer</a>	<span>✓ OK</span>	Node Port	34.118.234.148:27017 TCP	1/1	leaf-image-management-system	<a href="#">cluster-1</a>
<input type="checkbox"/> <a href="#">mongodb-producer</a>	<span>✓ OK</span>	Node Port	34.118.227.224:27017 TCP	1/1	leaf-image-management-system	<a href="#">cluster-1</a>
<input type="checkbox"/> <a href="#">users</a>	<span>✓ OK</span>	Node Port	34.118.226.66:5050 TCP	1/1	leaf-image-management-system	<a href="#">cluster-1</a>
<input type="checkbox"/> <a href="#">zookeeper</a>	<span>✓ OK</span>	Node Port	34.118.233.103:2181 TCP	1/1	leaf-image-management-system	<a href="#">cluster-1</a>

## Ingress

Gateways, Services & Ingress						
Cluster		Namespace				
cluster-1 (us-east1-b)		leaf-image-managemen...		Reset	Save	<a href="#">Learn</a>
<a href="#">Gateways</a> <a href="#">Routes</a> <a href="#">Services</a> <a href="#">Policies</a> <a href="#">Ingress</a>						
<input type="checkbox"/> <a href="#">Name ↑</a>	Status	Type	Frontends	Services	Namespace	Clusters
<input type="checkbox"/> <a href="#">camera-ingress</a>	<span>✓ OK</span>	External HTTP(S) LB	34.120.253.19/	<a href="#">camera</a>	leaf-image-management-system	<a href="#">cluster-1</a>
<input type="checkbox"/> <a href="#">consumer-ingress</a>	<span>✓ OK</span>	External HTTP(S) LB	34.120.106.39/image-plant	<a href="#">consumer-api</a>	leaf-image-management-system	<a href="#">cluster-1</a>
<input type="checkbox"/> <a href="#">db-synchronizer-ingress</a>	<span>✓ OK</span>	External HTTP(S) LB	34.144.252.58/	34.144.252.58/metrics	<a href="#">db-...</a>	leaf-image-management-system
<input type="checkbox"/> <a href="#">image-api-ingress</a>	<span>✓ OK</span>	External HTTP(S) LB	34.8.14.150/ping	34.8.14.150/image-plant	34.8.14.150/docs	<a href="#">image-api</a>
<input type="checkbox"/> <a href="#">users-ingress</a>	<span>✓ OK</span>	External HTTP(S) LB	35.244.192.167/	<a href="#">users</a>	leaf-image-management-system	<a href="#">cluster-1</a>

## Utilization



## Cost estimation

The cluster, as it is, costs around \$305/month (with batching and streaming)

Estimated monthly cost [Preview](#)

\$304.11

That's about \$0.42 per hour

Pricing is based on the resources you use, management fees, discounts and credits.

[Learn more ↗](#)

On GKE, a batch solution that runs every 10 minutes could be cheaper than a continuously running event-driven architecture if scaling nodes only during batch execution (configure nodes' autoscaling), but if the batches run long enough to occupy the cluster most of the time in between runs, costs approach those of streaming.

For example, a 3-node cluster with `e2-custom-4-4096` nodes as above (~\$0.42/hr) would cost roughly \$3,680/year if running 24/7, whereas a batch job occupying ~10% of that time could cost ~\$368/year, while a streaming setup running continuously would be the full \$3,680/year;

It all depends on how “real-time” processing should be and how long it takes to process a batch.

Compared to purchasing own hardware, which might require \$5k–10k upfront plus \$1k–2k/year in electricity, cooling, and maintenance, GKE provides quite a benefit with auto-scaling, managed upgrades, and higher reliability, making it more flexible and often more cost-effective

for workloads with variable load, though dedicated hardware could be cheaper for consistently high utilization.